

Guided Linking: Dynamic Linking without the Costs

Motivation

Optimizations don't cross dynamic linking boundaries

```
// plugin.so      // library.so
void foo(T* p) {   int bar(T* p) {
    int n = bar(p); return p->size;
    // ...        }
}
```

We want to inline bar into foo—but how?

Guided Linking

- A new extension of LTO
- Works on existing, **unmodified software**
- Uses existing, **unmodified optimizations**
- Enables new deduplication technique
- Code: <https://github.com/yotann/bcdb>

Related Work

- Ziegler et al. 2019; Agadakos et al. 2019
 - Performs dead code elimination for shared libraries
 - Can't apply any other optimizations
- Dietz & Adve 2018 (Software Multiplexing)
 - Converts dynamic linking to static linking
 - Inflexible and incompatible

Authors



Sean Bartell

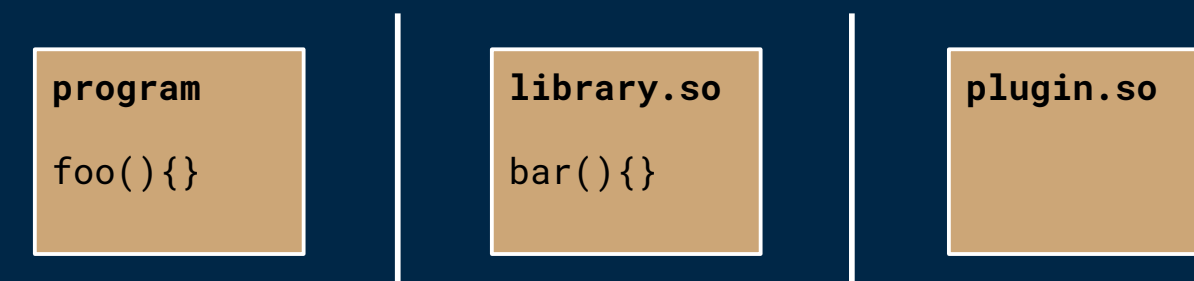
Will Dietz

Vikram S. Adve

University of Illinois


Problem #1

Compilers optimize each dynamic object **separately**




Solution

Developer provides **multiple objects** at once in the form of compiler IR



We move code to a **merged library**



And apply **existing optimizations** and **deduplication**

Optimized Set


- Arbitrary set of programs, libraries, plugins
- Must be compiled into IR form
- Must be distributed as a single unit
- Could be one package, an entire Docker container, ...

Problem #2

Dynamic linking is **unpredictable**

LD_DYNAMIC_WEAK? LD_PRELOAD?
Modified libraries? /etc/ld.so.cache?
Set-user-ID? LD_LIBRARY_PATH?
Interposing definitions?

We can maintain perfect compatibility using **indirection** through the dynamic linker, but this prevents optimizations




Solution

Developer provides **constraints**

e.g.: "This foo() will never be overridden by a different foo()."

Constraints let us **avoid indirection**





Constraints

Four constraints can be applied to each symbol:

- NoOverride:** definition will not be overridden by an external definition.
- NoUse:** definition will not be used externally or with dlsym().
- NoPlugin:** code will not be used in a plugin.
- NoWeak:** no weak uses or external definitions.

With all four, we can almost optimize the code as if it were statically linked.

Results



Shrinks & speeds up Clang+LLVM (depending on constraints)

